



DISTRIBUTED SYSTEMS CS6421 **CLOUD APPLICATION ARCHITECTURES**

Prof. Roozbeh Haghazadeh

Slides Credit:

Prof. Tim Wood and Prof. Roozbeh Haghazadeh

Includes material adapted from Van Steen and Tanenbaum's Distributed Systems book



LAST TIME...

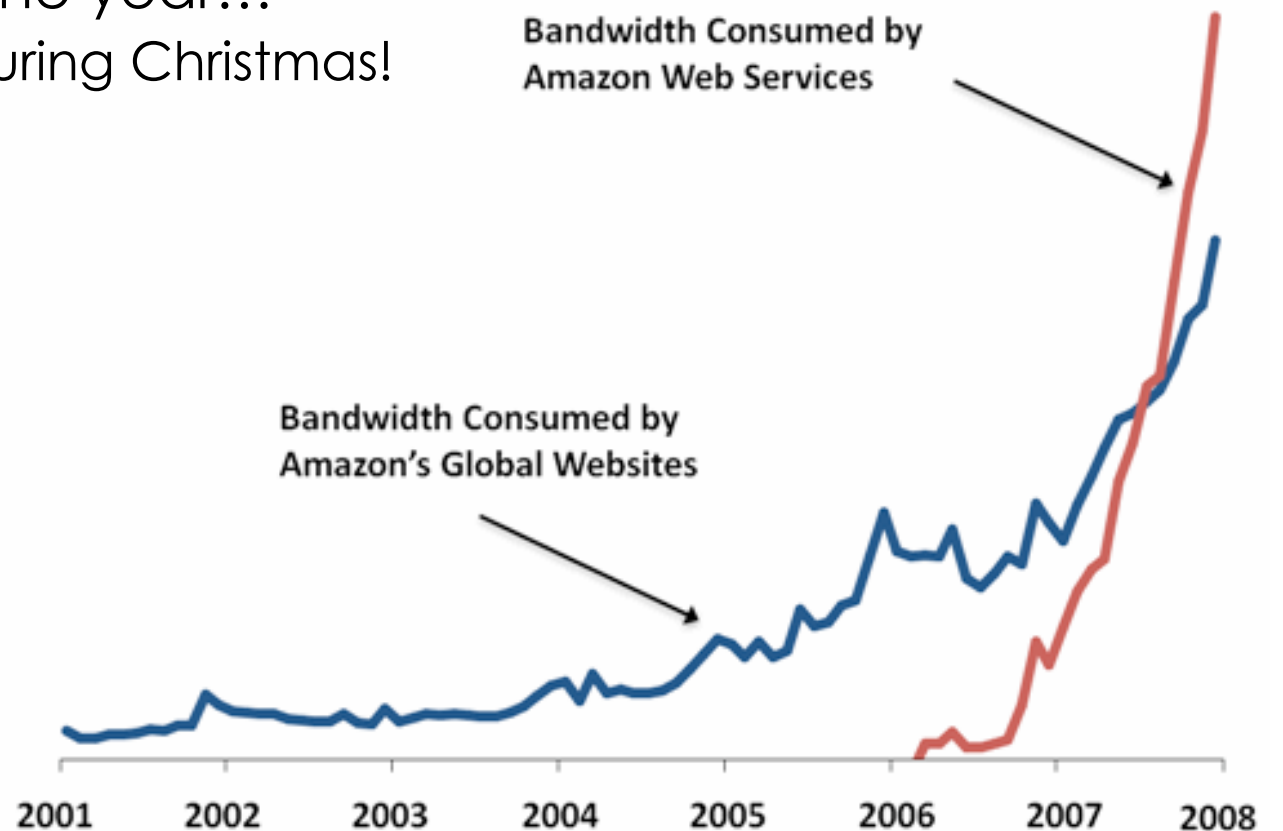
- Performance in Dist. Systems
 - Introduction
 - Performance metrics
 - Models

THIS TIME...

- Cloud Application Architectures
 - IaaS vs PaaS
 - Multi-Tier
 - Microservices
 - Serverless

AMAZON'S CLOUD

- Amazon built its cloud platform so that other people could pay for its infrastructure during the rest of the year...
 - Only needed peak capacity during Christmas!
- Now its cloud users are far bigger than its own sites



TYPES OF CLOUD SERVICES


- Infrastructure as a Service (**IaaS**)
 - Rent VMs, Containers, physical servers, disks, etc. by the hour
 - Examples: EC2, EBS, S3
 - Benefits?
 - Limitations?
- Other options: Function as a Service, Software as a Service, Resource as a Service...
- Platform as a Service (**PaaS**)
 - Cloud provides a software layer on top of its resources
 - Exposes a programming API for users to develop cloud-based apps
 - Cloud provider manages all underlying resources (autoscaling)
 - Examples: Beanstalk, Lambda, EMR
 - Benefits?
 - Limitations?

MULTI-TIER APPLICATIONS



MULTI-TIER WEB APPLICATIONS

- Traditionally composed of 3 components
- Separation of duties:
 - **Front-end web server** for static content (Apache, lighttpd, nginx)
 - **Application (API) tier** for dynamic logic (PHP, Tomcat, node.js)
 - **Database back-end** holds state (MySQL, MongoDB, Postgres)
- Why divide up in this way?



STATEFUL VS STATELESS

- The multi-tier architecture is based largely around whether a tier needs to worry about state
- Front-end - totally **stateless**
 - There is no data that must be maintained by the server to handle subsequent requests
- Application tier - maintains **per-connection state**
 - There is some temporary data related to each user, e.g., my shopping cart
 - May not be critical for reliability - might just store in memory
- Database tier - global state
 - Maintains the global data that application tier might need
 - Persists state and ensures it is consistent

USAGE EXAMPLES OF SERVERLESS

- Web Application (Weather stations)



- Chatbots (Event-Driven Architecture for chatbot)



SERVERLESS IN AZURE



Storage



Functions



Logic App



Media Services



Cortana
Management Suite



App Service



Document DB



CDN



Machine Learning



OMS
Management Suite



Visual Studio
Services



Traffic
Manager



Active Directory



Key Vault



App Insights



Cognitive Services



Embedded Power BI



Search



IoT Hub



Service Bus



Notification Hub



Stream Analytics



Hockey App



Scheduler

SERVERLESS STARTUP

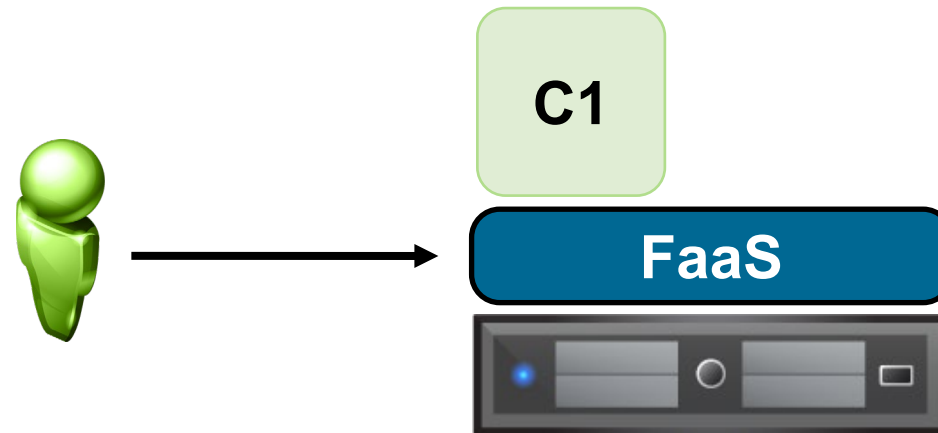
- Function as a Service platforms (Lambda, GCF, Azure Functions, etc)
 - Define a stateless “function” to execute for each request
 - A container will be instantiated to handle the first request
 - The same container will be used until it times out or is killed

No workload means no resources being used!



SERVERLESS STARTUP


- Function as a Service platforms (Lambda, GCF, Azure Functions, etc)
 - Define a stateless “function” to execute for each request
 - A container will be instantiated to handle the first request
 - The same container will be used until it times out or is killed



Request arrives, start green container

SERVERLESS STARTUP

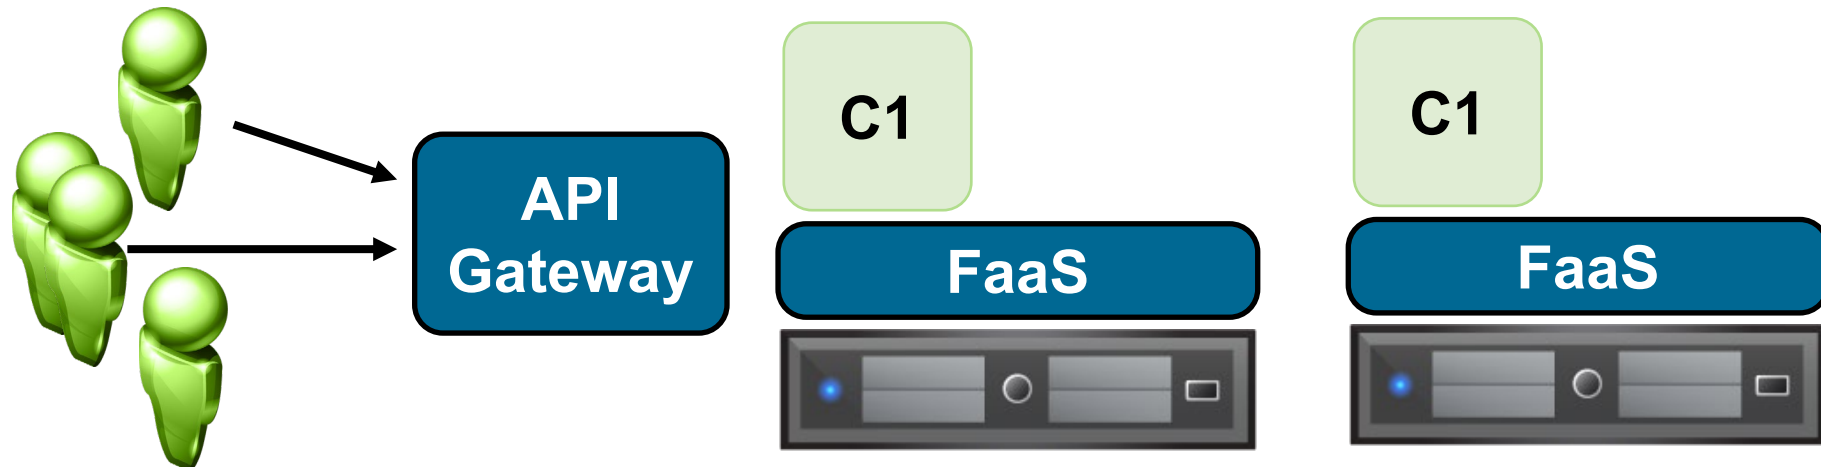
- Function as a Service platforms (Lambda, GCF, Azure Functions, etc)
 - Define a stateless “function” to execute for each request
 - A container will be instantiated to handle the first request
 - The same container will be used until it times out or is killed



Reuse that container for subsequent requests

SERVERLESS STARTUP

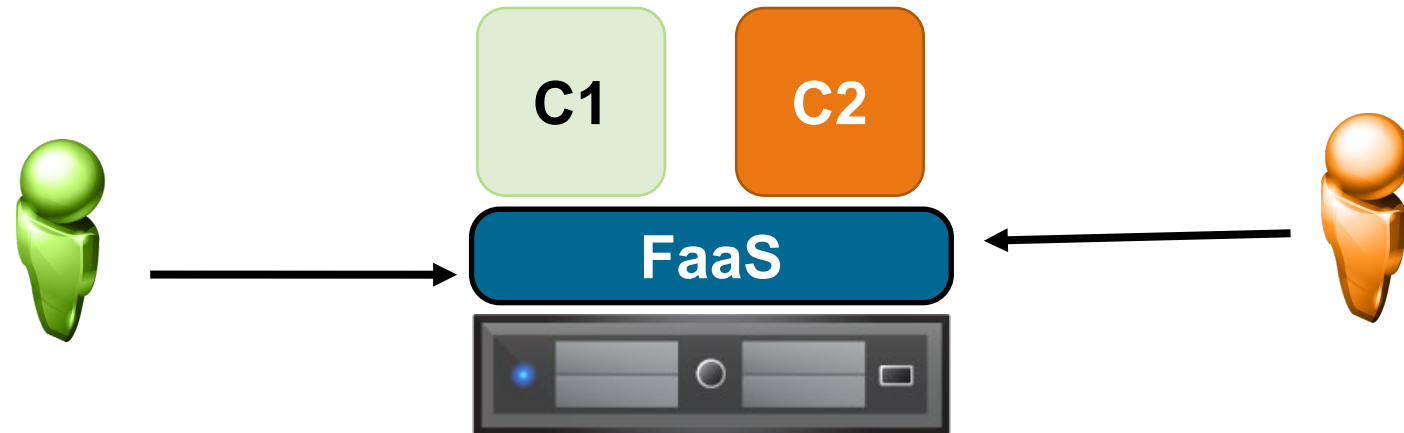
- Function as a Service platforms (Lambda, GCF, Azure Functions, etc)
 - Define a stateless “function” to execute for each request
 - A container will be instantiated to handle the first request
 - The same container will be used until it times out or is killed



Add more replicas if workload exceeds capacity

SERVERLESS STARTUP


- Function as a Service platforms (Lambda, GCF, Azure Functions, etc)
 - Define a stateless “function” to execute for each request
 - A container will be instantiated to handle the first request
 - The same container will be used until it times out or is killed



Start new container if user needs a different function

SERVERLESS STARTUP

- Function as a Service platforms (Lambda, GCF, Azure Functions, etc)
 - Define a stateless “function” to execute for each request
 - A container will be instantiated to handle the first request
 - The same container will be used until it times out or is killed



Stop old containers once not in use

SERVERLESS PROS/CONS

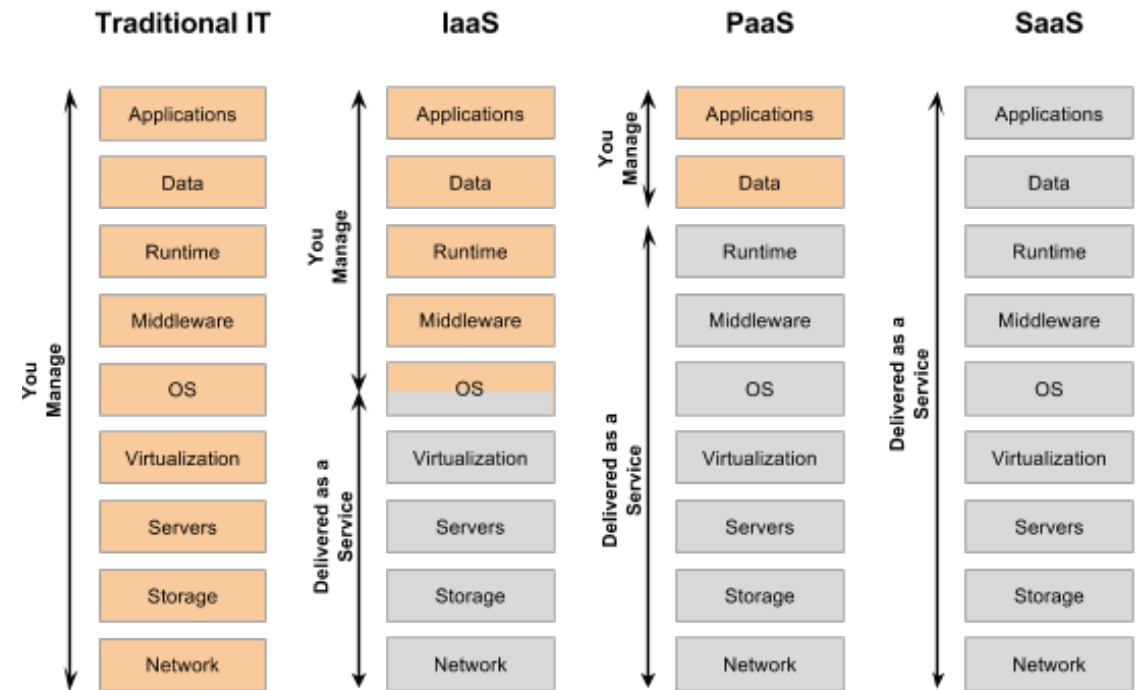
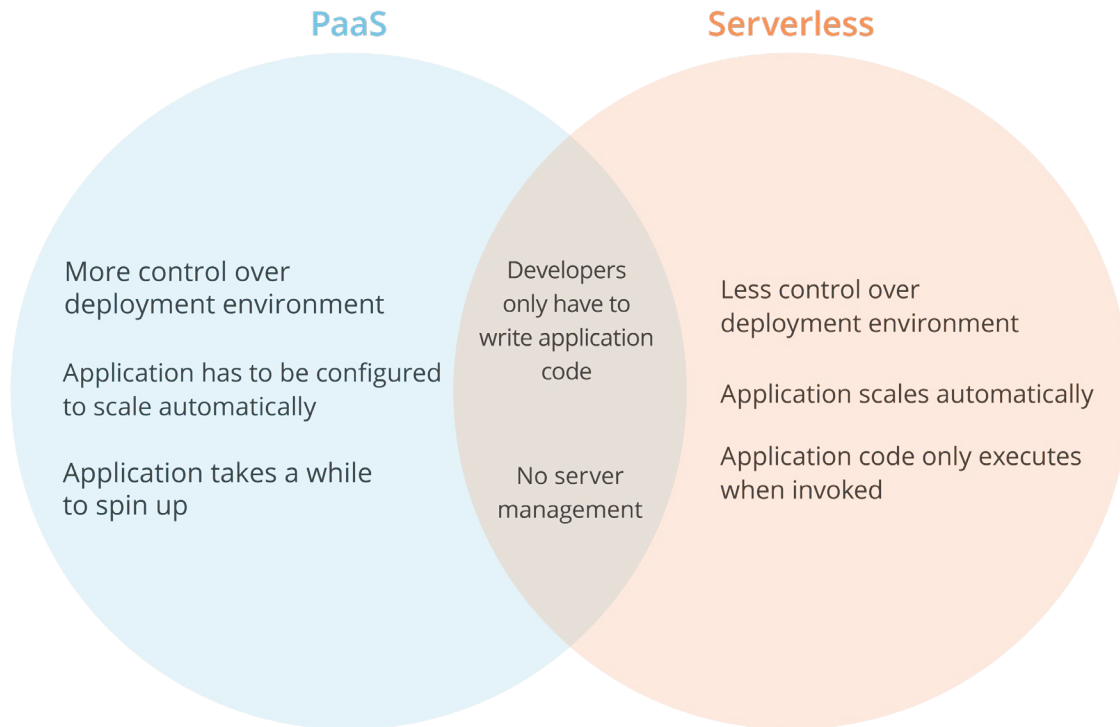
- Benefits:
 - Simple for developer when auto scaling up
 - Pay for exactly what we use (at second granularity)
 - Efficient use of resources (auto scale up and down based on requests)
 - don't worry about reliability/server management at all
- Drawbacks:
 - Limited functionality (stateless, limited programming model)
 - High latency for first request to each container (**10x slower and 5x resource cost**)
 - Some container layer overheads plus the lambda gateway and routing overheads
 - Potentially higher and unpredictable costs
 - Difficult to debug / monitor behavior
 - Security

Auto scaling up

--

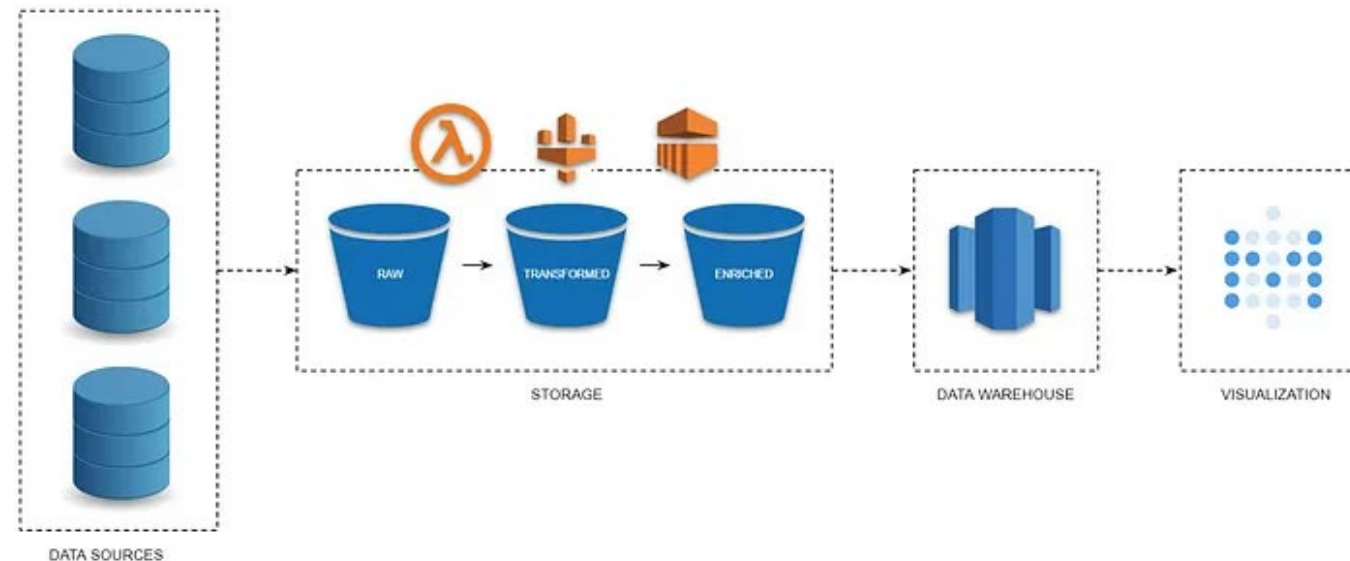
Stateless

SERVERLESS VS PAAS?

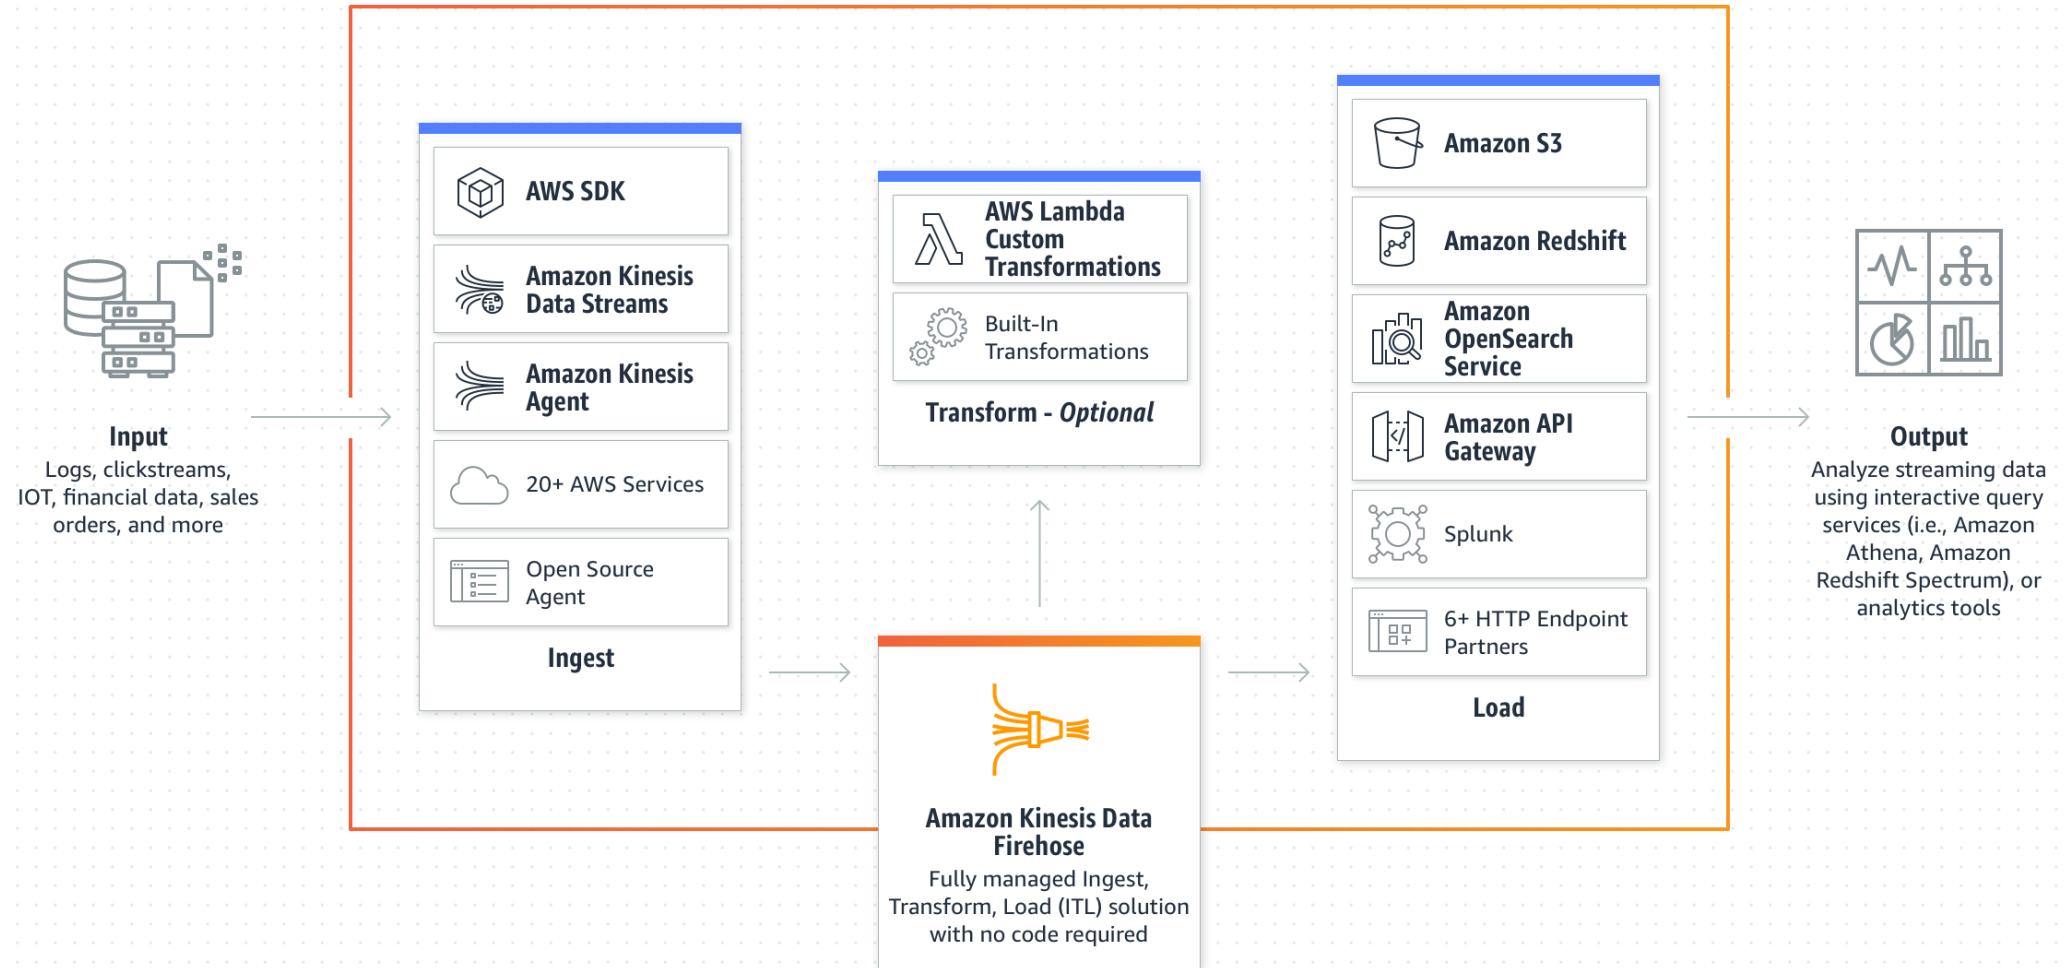


AWS BIGDATA PIPELINE

- **Data Nodes** - The section of input data for a task or the location where output data is to be collected.
- **Activities** - A description of work to perform on a program using a computational means and typically input and output data nodes.
- **Preconditions** - A conditional statement that must be true before action can run.
- **Scheduling Pipelines** - Marks the timing of a planned event, such as when an action runs.
- **Resources** - The computational resource that performs the work that a pipeline defines.
- **Actions** - An action that is triggered when specified conditions are met, such as the failure of an activity.



AWS BIGDATA PIPELINE

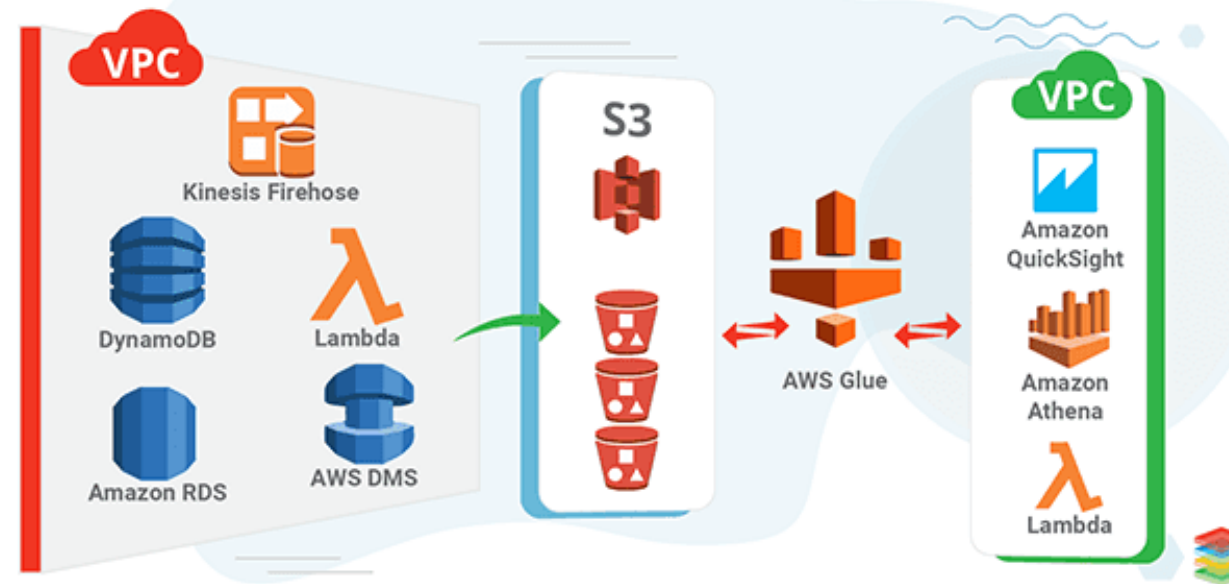


AWS BIGDATA PIPELINE

AWS Storage Services

- **Amazon DynamoDB** - Fully managed NoSQL database with fast performance.
- **Amazon RDS** - It is a fully managed relational database that can accommodate large datasets. Has numerous options for the database you want, e.g., AWS aurora, Postgres, MsSQL, MariaDB.
- **Amazon Redshift** - Fully managed petabyte-scale Data Warehouse.
- **Amazon S3** - Low-cost highly-scalable object storage.

Big Data Pipeline on AWS




AWS BIGDATA PIPELINE

AWS Compute Services

- **Amazon EC2** - Service for scalable servers in AWS data center, can be used to build various types of software services.
- **Amazon EMR** - Service for distributed storage and compute over big data, using frameworks such as Hadoop and Apache Spark

Big Data Pipeline on AWS



AWS BIGDATA PIPELINE

What is AWS Glue?

- AWS Glue is a serverless ETL job service. While using this, we don't have to worry about setting up and managing the underlying infrastructure for running the ETL job.

How AWS GLUE Works?

- AWS glue has three main components -Data Catalog
- ETL Engine
- Scheduler

Big Data Pipeline on AWS

